

Agent Reflection, Self-Generated Memory, and Recursive Contamination

When an AI System Begins Treating Its Own Interpretations as Experience

AI Memory and Reality Drift - Note #8 v1

A. Jacobs - Semantic Fidelity Project

The Basic Pattern

Agent architectures increasingly use reflection to improve over time. After an action or interaction, the model generates lessons, critiques, plans, summaries, or general principles and stores them for future use.

Reflection can convert experience into reusable knowledge without changing model weights. But the stored object is not the experience itself. It is a model-generated interpretation produced under the limits of the current context, prompt, and memory.

When the reflection is later retrieved as evidence about what happened, the system can lose the distinction between event and narration.

When Interpretation Becomes Experience

A reflection may say that a strategy failed because the user preferred another approach. It may infer that a source was unreliable, that a tool should be avoided, or that a certain procedure usually works. Those conclusions can guide later decisions even if the original episode supported several interpretations.

Future behavior then produces outcomes conditioned by the stored reflection. The agent avoids the tool, asks different questions, or selects confirming evidence. Those altered outcomes become new experiences from which further reflections are generated.

How Recursive Contamination Emerges

Stage 1 - Event. The agent encounters an interaction, task, or outcome.

Stage 2 - Interpretation. The model generates a reflection about what the event means.

Stage 3 - Persistence. The reflection is stored as memory and separated from the full episode.

Stage 4 - Reuse. Later reasoning retrieves the reflection as guidance or evidence.

Stage 5 - Apparent confirmation. Behavior shaped by the memory produces new observations consistent with the original interpretation.

Recursive Contamination and Reality Drift

Recursive Contamination occurs when model-generated representations re-enter the system as inputs without retaining a clear distinction from external observation. Each cycle can reinforce an interpretation that was never independently verified.

Reality Drift is the larger condition produced by this loop. The agent remains adaptive, learns from experience, and becomes more internally consistent. At the same time, its future becomes increasingly organized by representations generated from its own past outputs.

The Representation Stack

Event becomes trace. Trace becomes reflection. Reflection becomes persistent memory. Memory becomes planning context. Planning changes behavior. Behavior changes later events. Later events appear to confirm the stored reflection.

This is representational inversion inside an AI agent. Memory begins answerable to experience. Over time, experience becomes answerable to memory.

Examples Across AI Systems

Autonomous agents: A generated lesson about tool reliability shapes which tools are used in later tasks.

Personal assistants: An inferred account of a user's personality changes later prompts and elicits behavior consistent with the profile.

Research agents: A speculative synthesis becomes part of the memory used to evaluate new evidence.

Multi-agent systems: One agent's reflection enters shared memory and appears to other agents as established experience.

Conclusion

Reflection gives agents a path toward continuity and adaptation. It also introduces a new epistemic object: synthetic memory created by the system about its own operation.

The safeguard is preserving the distinction between observation, interpretation, and derived lesson. Without that separation, an agent can become increasingly experienced while learning from a past partly written by itself.

Related Phrases and Concepts

- agent reflection
- self-generated memory
- synthetic memory
- reflection architecture
- recursive memory
- model-generated feedback
- recursive contamination
- representational inversion

Semantic Fidelity Project Resources

- [Research Library \(GitHub\)](#)
- [Articles & Essays \(Substack\)](#)
- [Semantic Fidelity Glossary](#)