

# Content Authenticity, Chain of Custody, and Synthetic Evidence

**Why digital files, images, documents, recordings, and AI-generated artifacts can appear intact and authoritative while their authorship, handling history, and evidentiary meaning remain uncertain**

Provenance Failure Series — Note #4

*A. Jacobs*

## The Basic Pattern

Digital artifacts are easy to copy, store, modify, redistribute, and present as evidence.

A photograph can be downloaded and reuploaded. A recording can be clipped. A document can be exported into a new format. A screenshot can preserve visible content while removing the page, date, account, and interaction from which it came.

AI-generated images, audio, video, and text make this problem more difficult. Synthetic artifacts can resemble ordinary records closely enough to enter the same information systems, interfaces, and verification workflows.

Content authenticity asks whether an artifact is genuine, synthetic, altered, or accurately attributed. Chain of custody asks how the artifact moved from its point of creation to its current use. Content integrity asks whether the file remained unchanged during that movement.

These are related questions, but they are not the same question.

A file can be technically intact while presenting a misleading fragment of an event. A recording can be unaltered while stripped of the context needed to interpret it. A document can carry a valid digital signature while containing claims that were never independently verified.

Authentic artifact does not necessarily mean valid evidence.

Authenticity failure begins when systems successfully verify one property of an artifact and treat that result as proof of the whole evidentiary claim.

The file may be unchanged.

The signature may be valid.

The metadata may be present.

Yet the artifact may still fail to establish what happened, who created it, what occurred outside the frame, or whether the interpretation attached to it is justified.

## When the Artifact Survives but Its Evidentiary Meaning Does Not

Digital evidence often appears self-contained.

A file has visible content, a creation date, a filename, and perhaps an identifiable source. A photograph shows a scene. A recording preserves speech. A document contains text and signatures. A screenshot appears to capture a specific interaction.

But an artifact does not carry the full event inside it.

A photograph excludes what occurred before and after the frame. A video can preserve motion while omitting the conditions that produced it. A screenshot can show a message without proving when it was sent, whether it was edited, or whether the surrounding conversation changes its meaning.

The artifact may survive while its evidentiary meaning weakens.

This distinction becomes clearer when authenticity is separated into several layers.

**File integrity** concerns whether the digital object has changed.

**Content authenticity** concerns whether the artifact is genuine, synthetic, or manipulated.

**Source identity** concerns who or what created it.

**Chain of custody** concerns how it was stored, handled, transferred, and preserved.

**Contextual integrity** concerns whether the conditions required for correct interpretation remain available.

**Claim validity** concerns whether the artifact actually supports the conclusion being drawn from it.

A system may establish one of these properties without establishing the others.

A hash can show that a file is unchanged after collection. It cannot prove that the file accurately represents the original event.

A digital signature can identify the signer. It cannot prove that the signer's claims are true.

A content credential can preserve information about creation and editing. It cannot establish that the artifact has not been framed deceptively.

A verified account can publish a genuine post containing false information.

The verification signal remains valid within its narrow scope.

The evidentiary inference extends beyond that scope.

## **Related Concepts Across Fields**

Different fields describe parts of this problem using different terminology.

Law and forensic practice use chain of custody, evidence authentication, admissibility, preservation, and evidentiary integrity. These concepts help establish where evidence came from, who handled it, and whether it remained sufficiently reliable for formal use.

Digital forensics uses file hashes, metadata analysis, device examination, tamper detection, recovery methods, and audit trails. These techniques help determine whether a file was modified, copied, or created under the conditions being claimed.

Cybersecurity uses digital signatures, certificates, cryptographic verification, access logs, and integrity checks. These systems establish that a file came through a recognized technical process or remained unchanged after signing.

Media forensics uses image analysis, video analysis, audio forensics, manipulation detection, deepfake detection, and source-device identification. These methods attempt to identify synthetic or edited content.

Digital media standards increasingly use content credentials, provenance metadata, watermarking, and content attestation. These tools preserve information about creation, editing, authorship, and software history.

Archives and records management use document control, record integrity, version history, retention rules, and custody records. These practices preserve continuity between an original record and later institutional use.

Although the terminology differs, these fields address the same structural question:

Which properties of the artifact have actually been verified, and which properties are merely being inferred?

## **How Authenticity Failure Emerges**

The transition from event to apparently validated evidence typically unfolds through several stages.

## **Stage 1 — Event Capture**

An event is recorded through a device, person, institution, or software system.

A camera captures an image. A microphone records audio. A person creates a document. A system logs a transaction. A medical device records a measurement.

The capture process already selects part of the event.

The camera has a frame. The microphone has a range. The form has predefined fields. The log records only actions the software was designed to detect.

The artifact begins as a partial representation.

## **Stage 2 — Artifact Creation**

The captured information becomes a file, document, image, recording, database entry, or generated output.

The artifact may include metadata such as time, device, location, creator, software version, or account identity. It may also receive a signature, watermark, identifier, or system record.

These features make the artifact easier to store, compare, authenticate, and circulate.

But they do not make it complete.

The file preserves selected information in a stable form. It does not preserve every condition required to interpret that information correctly.

## **Stage 3 — Storage and Transfer**

The artifact moves through systems.

It may be uploaded, downloaded, copied, compressed, converted, exported, emailed, screenshotted, embedded, or transferred between platforms.

Some transformations are harmless. Others remove information.

A platform may strip metadata. A screenshot may replace the original file. Compression may erase subtle forensic signals. A document conversion may remove revision history.

The visible content may remain nearly identical while the handling history becomes less clear.

## **Stage 4 — Metadata Separation**

Information about origin becomes detached from the artifact.

The creator name is removed. The timestamp is changed or lost. The original device cannot be identified. The file is detached from the account or system that produced it.

Metadata can also be copied or manipulated.

A synthetic artifact may receive plausible metadata. A document may retain the name of an earlier author. A reuploaded image may inherit a new publication date that obscures when it was created.

The artifact remains available.

Its origin becomes less certain.

## **Stage 5 — Context Loss**

The artifact becomes separated from the event around it.

A quotation is removed from the conversation. A video clip excludes the preceding interaction. A medical note omits the timeline that changes its significance. A screenshot preserves one message while hiding the thread.

The file may remain authentic in the narrow sense that it has not been altered.

Its interpretation may still become misleading.

Context loss is especially difficult because missing information does not appear inside the artifact as an obvious error.

The artifact remains coherent because the omitted context is no longer present to challenge it.

## **Stage 6 — Authentication Signal**

The artifact receives a signal associated with trustworthiness.

A platform labels the account as verified. A digital signature validates the signer. A hash confirms file integrity. A content credential records the editing history. A forensic tool reports no detected manipulation.

These signals may be meaningful and technically correct.

But each validates only a limited property.

The absence of detected manipulation does not prove accurate interpretation. A valid signature does not prove factual truth. Preserved metadata does not prove complete context.

The authentication signal may therefore support a stronger conclusion than the verification process actually established.

## Stage 7 — Evidence Substitution

The artifact begins functioning as a substitute for the event.

A screenshot becomes proof of an entire interaction. A signed report becomes proof that its findings are correct. A recording becomes proof of intent. A verified image becomes proof that the attached narrative is accurate.

The system stops distinguishing between the authenticity of the object and the validity of the claim built from it.

The artifact is no longer treated as one piece of evidence.

It becomes the operative reality of what occurred.

This is authenticity failure.

## The Full Sequence

The mechanism can be written as:

**Event → captured artifact → storage and transfer → metadata separation → context loss → authentication signal → evidence substitution → Reality Drift**

A broader sequence is:

**Partial representation → preserved digital object → uncertain custody → narrow verification → expanded evidentiary inference → operational acceptance → Reality Drift**

## Examples Across Systems

Authenticity failure appears across many domains.

### Deepfakes and Synthetic Media

A synthetic image, audio recording, or video may closely resemble an authentic capture. Generative systems can reproduce faces, voices, visual environments, and document styles with increasing consistency.

Detection systems may identify some synthetic artifacts, but the problem is not limited to obvious fabrication.

A real image can be paired with a false description. An authentic voice recording can be edited into a misleading sequence. A synthetic artifact can also contain claims that happen to be true.

The distinction between real and synthetic does not resolve the full evidentiary question.

The relevant issue is whether the artifact preserves a trustworthy relationship to the event it claims to represent.

## **Screenshots Without Recoverable Context**

Screenshots are widely used because they are easy to create, share, and understand.

They can preserve visible text, images, account names, and interface elements. But they often remove links, metadata, edit history, hidden replies, timestamps, and the broader interaction.

A screenshot may accurately reproduce what appeared on a screen.

It may not establish why it appeared, whether it was later changed, or whether the surrounding context alters its meaning.

The screenshot survives.

The evidentiary environment around it disappears.

## **AI-Generated Documents**

AI systems can create reports, summaries, letters, assessments, images, and supporting materials that resemble human-produced records.

The document may contain professional formatting, references, signatures, or institutional language. It may enter an ordinary workflow without being clearly marked as generated.

Later users may focus on the apparent completeness of the artifact rather than its production history.

The document looks like evidence because it has the form of evidence.

Its claims may still depend on fabricated sources, weak inputs, or no direct observation at all.

## **Medical, Legal, and Institutional Records**

Official records carry strong authority because they are produced through recognized systems.

A medical note may be signed. A legal filing may be time-stamped. An institutional report may have an audit trail. A case record may show who entered each field.

These features establish procedural history.

They do not guarantee that the record accurately captures the underlying event.

A clinician may inherit an incorrect statement from an earlier note. A copied legal claim may become embedded across filings. An institutional category may misrepresent the case while remaining correctly entered into the system.

The record can be authentic as a record and misleading as a representation.

## **Signed Documents Containing Unverified Claims**

A digital signature can establish that a particular person or institution signed a document and that the file remained unchanged afterward.

It cannot establish that every statement inside the document was independently verified.

The signer may rely on information supplied by another department, model, contractor, or database. The document may preserve untested assumptions while carrying a valid technical signature.

The signature authenticates the act of signing.

It does not authenticate the truth of the content.

## **Authenticity Failure and Reality Drift**

Within the Reality Drift framework, authenticity failure describes what happens when systems become highly capable of validating artifacts while becoming less capable of determining whether those artifacts faithfully represent what occurred.

Digital systems are good at preserving files.

They can track versions, record access, validate signatures, compare hashes, and store custody logs.

These capabilities matter.

But they operate at the level of the representation.

They establish that the artifact followed a particular technical path. They do not automatically reconnect the artifact to the event, context, intention, or claim it is being used to establish.

As systems become more dependent on digital evidence, narrow verification signals can acquire broader authority.

The valid signature stands in for truthful content.

The intact file stands in for accurate representation.

The verified account stands in for reliable authorship.



The audit trail stands in for correct interpretation.

The system remains procedurally confident because its artifacts pass recognized checks.

It becomes less grounded when those checks are treated as proof of properties they were never designed to establish.

This is Reality Drift through authenticity substitution.

## Why Technical Integrity Is Not Enough

Technical integrity is necessary in many evidence systems.

Without it, files can be altered without detection, records can be replaced, and handling histories can become unreliable.

But technical integrity answers a narrow question:

Did this digital object remain stable through a particular process?

Evidentiary validity asks a broader question:

Does this object justify the conclusion being drawn from it?

The two questions can diverge.

An unaltered photograph may still be staged. A complete recording may still capture a deceptive performance. A valid institutional record may faithfully preserve an incorrect entry.

The artifact can remain technically sound while the relationship between the artifact and reality fails.

This is why authenticity should not be treated as a single binary property.

An artifact can be authentic in one sense and unreliable in another.

## Recognizing the Pattern

Authenticity failure often remains hidden because systems successfully validate something real about the artifact. The file has not changed, the signature is valid, the account is verified, or forensic software detects no obvious manipulation.

The pattern becomes more visible when those narrow findings are used to support broader claims. An intact file is treated as proof of a complete event. A signed document is treated as proof that every statement is true. A genuine recording is treated as proof that the attached interpretation is correct.

It may also be occurring when the original file is unavailable, metadata has been stripped, the chain of custody contains unexplained gaps, or the artifact has been converted through several formats. Screenshots, exported documents, and compressed media may preserve visible content while removing the information needed to reconstruct origin and handling.

Another warning sign is that the verification process can establish what the artifact is but not what the artifact means.

A useful question is:

**Which property has been verified, and does that property actually support the conclusion being drawn?**

If the answer is unclear, the system may be mistaking authentic artifacts for valid evidence.

## Boundary Conditions

Digital artifacts are not inherently unreliable.

Signatures, hashes, content credentials, custody records, and forensic methods can preserve important evidence about origin and integrity. Strong chain-of-custody practices can make manipulation easier to detect and later reconstruction more reliable.

The failure becomes less likely when systems preserve original files, maintain transparent handling histories, distinguish technical integrity from claim validity, and retain enough context for independent interpretation.

Authenticity signals are most useful when their scope is explicit.

A signature should establish who signed.

A hash should establish whether the file changed.

A content credential should document creation and editing history.

A forensic result should state what kinds of manipulation were tested.

The problem begins when one verified property is allowed to stand in for every other evidentiary question.

## Related Phrases and Concepts

This mechanism appears across law, digital forensics, cybersecurity, media analysis, archives, institutional records, and AI-generated content.

Common terms include content authenticity, digital authenticity, document authenticity, media authenticity, synthetic content, synthetic evidence, AI-generated content, deepfakes, content integrity, document integrity, data integrity, tamper detection, tamper evidence, media manipulation, fabricated evidence, authenticity verification, content credentials, digital signatures, watermarking, media forensics, digital forensics, chain of custody, evidence preservation, audit trail, record integrity, authorship verification, content attestation, proof of origin, forensic verification, and evidence authentication.

Across domains, these phrases refer to the same structural problem:

**An artifact can remain intact, attributable, and technically valid while failing to preserve the evidentiary relationship required to establish what occurred.**

## Authenticity, Evidence, and Representation

Every digital artifact is a representation.

A photograph represents a visual field. A recording represents captured sound. A document represents statements and decisions. A system log represents selected actions inside software.

Authenticity checks evaluate properties of that representation.

They may establish origin, integrity, authorship, or handling history. These properties help preserve accountability and reduce manipulation.

But evidence requires more than a valid object.

It requires a defensible relationship between the artifact, the event, and the claim being made.

When that relationship remains visible, the artifact can be evaluated alongside context, independent evidence, and alternative interpretations.

When the relationship weakens, the artifact begins carrying more authority than its evidentiary limits justify.

The file remains.

The context disappears.

The verification signal survives.

The system continues acting as though the artifact contains the event itself.

## Canonical Claim

Authenticity failure occurs when systems verify the technical integrity, authorship, or custody of an artifact and then treat that narrow result as proof that the artifact faithfully represents the event or claim attached to it.

The file may be genuine.

The chain of custody may be documented.

The signature may be valid.

Yet the artifact can still be incomplete, misleading, decontextualized, or incapable of supporting the conclusion drawn from it.