

Provenance Loss, Source Traceability, and Reality Drift

Why information becomes harder to trace as summaries, copies, generated outputs, and derivative sources move farther from the original evidence

Provenance Failure Series — Note #1

A. Jacobs

The Basic Pattern

Information rarely moves directly from an original event to a final decision. An observation becomes a record, the record becomes a claim, and the claim may later enter a report, summary, database, search result, or AI-generated answer.

At each stage, the information may remain useful. But the connection to its origin can weaken.

Provenance describes the history of where information came from, how it was transformed, and which sources it depends on. Source traceability describes the ability to follow that history back toward the original evidence.

A document may have an author. A dataset may have a collection method. A photograph may contain a time, location, and device history. A claim may refer to an interview, experiment, observation, or institutional record.

These details establish where the information came from.

But information is rarely reused in its original form. A passage is quoted without its surrounding argument. A finding is reduced to a headline. A report becomes an executive summary. A chart is copied into a presentation. An AI-generated answer combines several sources into a single response.

Each transformation creates a new representation. The new version may preserve the central claim while removing part of the history required to evaluate it.

The author may disappear. The date may disappear. The method may disappear. The distinction between direct evidence and later interpretation may disappear.

Provenance loss begins when information continues circulating while the path back to its source becomes incomplete, uncertain, or impossible to reconstruct.

The claim remains available.

The origin becomes less visible.

The representation survives while its derivation history disappears.

When Information Becomes Detached From Origin

Source information is often preserved near the beginning of a process. The original record may identify who created it, when it was produced, how the evidence was collected, and what limitations shaped its meaning.

Those details become easier to lose as the information moves.

A quoted sentence may no longer include the qualification that followed it. A statistical result may circulate without the method used to produce it. A screenshot may preserve visible content while removing the surrounding page, date, link, and author. A generated summary may preserve the conclusion while obscuring which source supported each claim.

The information becomes easier to circulate than to trace.

The derivative version may also become more visible than the original. A summary may rank higher in search. A copied claim may appear across more websites. A generated answer may be encountered without the documents from which it was produced.

Readers then interact primarily with later representations.

The source remains technically available, but it no longer governs how the information is understood.

At this point, the representation begins functioning as the source.

The claim continues moving through the system. The provenance chain no longer moves with it.

Related Concepts Across Fields

Different disciplines describe parts of this pattern using different terminology.

Data management refers to data provenance and data lineage. These terms describe where data came from, how it moved through a system, and which transformations produced its current form.

Archives and records management use concepts such as chain of custody, record history, document traceability, and evidence preservation. These practices help establish who created a record, who handled it, and whether its integrity was maintained.

Journalism and scholarship rely on source attribution, citation history, primary sources, reference verification, and bibliographic records. These practices allow readers to distinguish direct evidence from later reporting, commentary, and interpretation.

Digital media systems increasingly use content credentials, authorship metadata, digital signatures, watermarking, and origin tracking. These tools attempt to preserve information about who created an artifact and whether it has been altered.

Machine learning systems use dataset lineage, model provenance, training data documentation, pipeline traceability, and data versioning. These records help explain which data, models, filters, and transformations contributed to an output.

Although the language differs, these concepts address the same structural question:

Can the current representation still be connected to the evidence from which it was derived?

How Provenance Loss Emerges

The transition from source evidence to provenance failure typically unfolds through several stages.

Stage 1 — Source Event

An event, observation, document, measurement, or experience produces the original evidence.

A researcher records a result. A journalist conducts an interview. A sensor captures a measurement. An employee documents an incident. A person takes a photograph.

At this stage, the information remains relatively close to the event that produced it. The original context may still be available.

Stage 2 — Initial Recording

The event is converted into a record.

The record may take the form of notes, data, audio, video, a database entry, a report, or a digital file. This makes the event easier to preserve, compare, and share.

But the recording is already selective.

A form captures only the fields it contains. A photograph captures only one angle. A transcript removes tone and setting. A dataset preserves only the variables that were measured.

The record is not the event itself. It is the first representation of the event.

Stage 3 — Extraction

A smaller claim is taken from the record.

A sentence is quoted. A number is selected. A conclusion is pulled from a report. A finding is separated from the method that produced it.

Extraction makes information easier to reuse. It also weakens context.

The extracted claim may no longer contain the limitations, uncertainty, or conditions that shaped its original meaning.

Stage 4 — Compression

The extracted material is shortened, grouped, or rewritten.

A report becomes a summary. Several sources become an overview. A complex event becomes a headline. A long analysis becomes a bullet point.

Compression reduces the amount of information that must move through the system. It also removes details that may later be necessary for interpretation or verification.

The claim remains.

The route back to the claim's origin becomes thinner.

Stage 5 — Derivative Reuse

The compressed representation is copied into new documents and systems.

A summary is quoted by another summary. A chart is reused without its original report. A claim appears in an article that cites another article. An AI-generated answer becomes source material for a webpage.

The derivative version begins circulating independently.

Later users may not know that they are encountering a representation of a representation. They may treat the derivative version as though it were an independent source.

Stage 6 — Metadata Separation

Information about origin becomes separated from the content itself.

The author is omitted. The source link breaks. The citation points only to a secondary document. The publication date is lost. The dataset name survives while the collection method disappears.

The artifact remains usable, but its derivation history becomes incomplete.

Provenance is no longer attached to the information as it moves.

Stage 7 — Source Substitution

The derivative version becomes easier to access than the original source.

Readers encounter the summary instead of the report. They encounter the screenshot instead of the underlying record. They encounter the AI answer instead of the source documents. They encounter repeated claims instead of independent evidence.

The original evidence may still exist, but it becomes operationally secondary.

The representation begins functioning as the source.

Stage 8 — Origin Reconstruction Failure

The path back to the source can no longer be reliably recovered.

Links lead to copies. Citations form loops. Multiple documents repeat the same wording. Metadata is missing or contradictory. Apparent sources may all depend on one earlier source that is no longer accessible.

The system still contains the claim.

It no longer contains enough information to explain where the claim came from, how it changed, or why it should be believed.

This is provenance failure.

Examples Across Systems

Provenance loss appears across many domains.

AI-Generated Information

AI systems often combine information from multiple documents into a single answer. The output may be coherent, accurate, and useful.

But the connection between each claim and its supporting source may be weak, incomplete, or absent.

A later user may copy the answer into another document. That document may then be indexed, summarized, and retrieved by another system.

The generated representation begins circulating as independent information. Its machine-mediated origin becomes harder to detect.

Journalism and Online Publishing

A news report cites another report. A blog summarizes the news report. A social post repeats the summary. A search result extracts the social post.

Several visible sources now repeat the same claim.

But the entire chain may depend on one original report.

The appearance of source diversity increases while the amount of independent evidence remains unchanged.

Institutional Reporting

A frontline event becomes a case record. The case record enters a category. The category becomes a metric. The metric appears in a report. The report becomes an executive summary.

Senior decision-makers may act on the summary without access to the original event.

The institution retains the reported conclusion while losing the local context that produced it.

Research and Scholarship

A paper cites an earlier study. Later papers cite the secondary paper rather than the original study.

The claim becomes widely repeated. Over time, its wording may become stronger or more definite than the original evidence justified.

The citation survives.

The original qualification disappears.

Digital Images and Documents

A photograph is downloaded and reuploaded. The filename changes. Metadata is removed. The image is cropped, compressed, or placed inside a screenshot.

The visible artifact survives.

Its time, location, authorship, and handling history become uncertain.

The image may still appear authentic while its provenance becomes difficult or impossible to establish.

Provenance Loss and Reality Drift

Within the Reality Drift framework, provenance loss describes one of the ways representations become detached from the realities they were created to preserve.

A system begins with information derived from an event, source, or observation. That information is recorded, extracted, compressed, and reused.

Each transformation may preserve enough structure for the representation to remain useful. But the connection to the original evidence gradually weakens.

Documents remain available. Search systems return results. AI systems generate answers. Institutional reports support decisions.

The system continues to circulate information.

Yet fewer participants can reconstruct where that information came from, which transformations shaped it, or whether visible sources are genuinely independent.

The system remains informed at the level of representation.

It becomes less answerable to the source reality beneath those representations.

This is Reality Drift through provenance loss.

Recognizing the Pattern

Provenance failure often remains invisible because the information still appears complete. The claim is written clearly, the document has a title, the report includes references, the answer may contain links, and the image looks intact.

These signals create the appearance of traceability. But apparent traceability is not the same as recoverable provenance.

The pattern becomes more visible when multiple sources repeat identical wording, citations lead only to secondary material, or the original document can no longer be located. It also appears when metadata is missing, inconsistent, or detached from the artifact.

Provenance loss may be occurring when summaries are easier to find than source material, copied claims appear more often than independent evidence, or source links resolve to pages that cite one another. AI-generated material may also be reused without disclosure, allowing derivative content to circulate as though it were independent.

An artifact may retain its visible content while losing its creator, date, handling history, or route back to the original evidence.

A useful question is:

Can the current representation be traced back through every major transformation to the evidence that originally supported it?

If the answer is no, the information may remain usable while its provenance has already begun to fail.

Related Phrases and Concepts

This mechanism appears across information systems, archives, research, journalism, digital media, and machine learning.

Common terms include information provenance, data provenance, content provenance, source provenance, data lineage, information lineage, source traceability, origin tracking, derivation history, provenance metadata, chain of custody, source attribution, authorship provenance, evidence lineage, primary source verification, source reconstruction, origin reconstruction, provenance loss, provenance failure, source substitution, derivative evidence, untraceable source, and broken provenance chains.

Across domains, these phrases refer to the same structural problem:

Information continues circulating after the history required to evaluate its origin has weakened or disappeared.

Provenance and Representation

Every record is a representation of something that occurred before the record existed.

A dataset represents observations. A report represents recorded activity. A citation represents a relationship between a claim and a source. A summary represents a larger body of material. A generated answer represents patterns extracted from many previous representations.

Provenance preserves the path between these layers.

It allows later users to distinguish the source from the interpretation, the observation from the summary, and the original evidence from the derivative copy.

When provenance is preserved, representations remain corrigible by their sources.

When provenance is lost, the representation may continue to circulate without the evidence needed to challenge, verify, or reinterpret it.

The information remains present.

The source becomes absent.

The system continues to know what was said while losing the ability to reconstruct why it should be believed.

This version follows the reference document's balance of short semantic paragraphs, grouped examples, and occasional isolated pivots rather than separating every sentence into its own block.